




File I/O



Introduction

- When a process uses a file,
- First, find the file in the file system → `open()`
 - Next, read or write data from the file → `read()/write()`
 - Finally, stop to use the file → `close()`



File Descriptor

file descriptor

- all open files are referred to by file descriptors.
- how to obtain file descriptor
 - return value of `open()`, `creat()`
- when we want to read or write a file,
 - we identify the file with the file descriptor
- file descriptor is the [index of user file descriptor table](#)
- `STDIN_FILENO(0)`, `STDOUT_FILENO(1)`, `STDERR_FILENO(2)`
 - defined in `<unistd.h>`
- Range of file descriptor
 - 0 ~ `OPEN_MAX` (63 in early many systems)

open()

```
#include <fcntl.h>
```

```
int open(const char *pathname, int oflag, ... /* mode_t mode */);
```

Returns: file descriptor if OK, -1 on error

 open/create a file and return a file descriptor.

- What does “...” mean?

- Third argument is used only when a new file is created.

open()

pathname

- The name of the file to open or create

oflag

- Access mode (One of three constants must be specified.)
- Only one of the following three constants should be specified
 - O_RDONLY
 - Open for reading only
 - O_WRONLY
 - Open for writing only
 - O_RDWR
 - Open for reading and writing

open()

oflag(cont.)

- The followings are optional.
- O_CREAT
 - Create the file if it doesn't exist.
 - Requires a third argument, **mode**.
- O_EXCL
 - Generate an error if O_CREAT is also specified and the file already exists.
- O_APPEND
 - Append to the end of file on each write.

open()

oflag(cont.)

- O_TRUNC
 - If the file exists and if it is successfully opened for either write-only or read-write, truncate its length to 0.
- O_SYNC
 - Any writes on the resulting file descriptor will block the calling process until the data has been physically written to the underlying hardware

open()

mode

- specifies the permissions to use if a new file is created.
- should always be specified when **O_CREAT** is in the flags, and is ignored otherwise.

return value

- return **the new file descriptor**, or -1 if an error occurred.
 - **the lowest numbered unused descriptor**

open()

example

```
int fd;  
fd = open("/etc/passwd", O_RDONLY);  
fd = open("/etc/passwd", O_RDWR);  
  
fd = open("ap", O_RDWR | O_APPEND);  
fd = open("ap", O_RDWR | O_CREAT | O_EXCL, 0644);
```

creat()

```
#include <fcntl.h>
```

```
int creat(const char *pathname, mode_t mode);
```

Returns: file descriptor opened for write-only if OK, -1 on error

Create a new file

- It is equivalent to
 - `open (pathname, O_CREAT|O_WRONLY|O_TRUNC, mode);`
- Note that the file is opened **only for writing**.

close()

```
#include <unistd.h>
```

```
int close(int fildes);
```

Returns: 0 if OK, -1 on error

Close an open file

- When a process terminates, all of its open files are closed automatically by the kernel.
- ➔ Many program often do not explicitly close open files.

read()

```
#include <unistd.h>
```

```
ssize_t read(int filedes, void *buf, size_t nbytes);
```

Returns: number of bytes read, 0 if end of file, -1 on error

 read up to *nbytes* from *filedes* into the buffer starting at *buf*

- read() starts at the file's current offset.
- Before a successful return, the offset is incremented by the number of bytes actually read.

 return value

- On success, the number of bytes read is returned.
- 0 indicates end of file.
- On error, -1 is returned.

read()


- ❏ the number of bytes actually read may be less than the amount requested.
 - If **the end of regular file** is reached before the requested number of bytes has been read.
 - When reading from a **terminal** device, up to **one line** is read at a time.
 - When reading from a **network**, **buffering** within the network may cause less than the requested amount to be returned.
 - When reading from a **pipe**, if the pipe contains **fewer bytes** than requested, read will return only what is available.
 - ...

write()

```
#include <unistd.h>
```

```
ssize_t write(int filedes, const void *buf, size_t nbytes);
```

Returns: number of bytes written if OK, -1 on error

 writes up to *nbytes* to the file referenced by *filedes* from the buffer starting at *buf*

- write **start** at the **file's current offset**.
- If O_APPEND was specified when the file was opened,
 - The file's offset is set to the end of file before write.

 return value

- On success, the number of bytes written is returned.
- On error, -1 is returned.

read()/write()

example

```
#include <unistd.h>
#include <stdio.h>
#define BUFFSIZE 8192

int main(void)
{
    int n;
    char buf[BUFFSIZE];

    while ((n=read(STDIN_FILENO, buf, BUFFSIZE))>0)
        if (write(STDOUT_FILENO, buf, n)!=n)
            printf("write error\n");

    if (n<0)
        printf("read error\n");
    exit(0);
}
```

read()/write()

Running result)

```
$ ./a.out
hello, world.
hello, world.
Are you enjoying this class?
Are you enjoying this class?
Ctrl + D
$
```


lseek()

```
#include <unistd.h>
```

```
off_t lseek(int filedes, off_t offset, int whence);
```

Returns: new file offset if OK, -1 on error

- ❏ Explicitly repositions an open file's offset
 - The offset for regular files must be **non-negative**.

❏ return value

- success: the resulting offset location as measured in bytes from the beginning of the file
- error: -1

lseek()

whence

- SEEK_SET
 - The offset is set to offset bytes from the **beginning of the file**.
- SEEK_CUR
 - The offset is set to its **current location** plus offset bytes.
- SEEK_END
 - The offset is set to **the size of the file** plus offset bytes.

lseek()


example

```
off_t curpos;  
curpos = lseek(fd, 0, SEEK_CUR);           // get the current offset  
  
lseek(fd, 0, SEEK_SET);  
lseek(fd, 0, SEEK_END);  
lseek(fd, -10, SEEK_CUR);  
lseek(fd, 100, SEEK_END);
```

lseek()

❏ hole

- The file's offset can be greater than the file's size.
 - Next write to the file will **extend the file**.
- It means that **a hole** in file is created and is allowed.
- read from the data in hole **returns 0**.



lseek()

example

```
#include "apue.h"
#include <fcntl.h>

char  buf1[] = "abcdefghij";
char  buf2[] = "ABCDEFGHJIJ";

int
main(void)
{
    int    fd;

    if ((fd = creat("file.hole", FILE_MODE)) < 0)
        err_sys("creat error");
    /* FILE_MODE is defined as 644 in "apue.h". */
}
```

lseek()

example(cont.)

```
if (write(fd, buf1, 10) != 10)
    err_sys("buf1 write error");
/* offset now = 10 */

if (lseek(fd, 16384, SEEK_SET) == -1)
    err_sys("lseek error");
/* offset now = 16384 */

if (write(fd, buf2, 10) != 10)
    err_sys("buf2 write error");
/* offset now = 16394 */

exit(0);
}
```

lseek()

Running result)

```
$ ./a.out
$ ls -l file.hole           check its size
-rw-r--r-- 1 sar          16394 Nov 25 01:01 file.hole
$ od -c file.hole          let's look at the actual contents
0000000  a b c d e f g h i j \0 \0 \0 \0 \0 \0
0000020  \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0
*
0040000  A B C D E F G H I J
0040012
```

byte offset in octal

od utility: dump files in octal.
-c: print the contents as characters

lseek()

Are the disk blocks allocated for hole?

```
$ ls -ls file.hole file.nohole
 8 -rw-r--r-- 1 sar      16394 Nov 25 01:01 file.hole
20 -rw-r--r-- 1 sar      16394 Nov 25 01:03 file.nohole
```

- Compare the sizes of file.hole and file.nohole
 - file.hole: with hole
 - ➔ 8 blocks are allocated
 - file.nohole: a file of the same size, but without holes.
 - ➔ 20 blocks are allocated

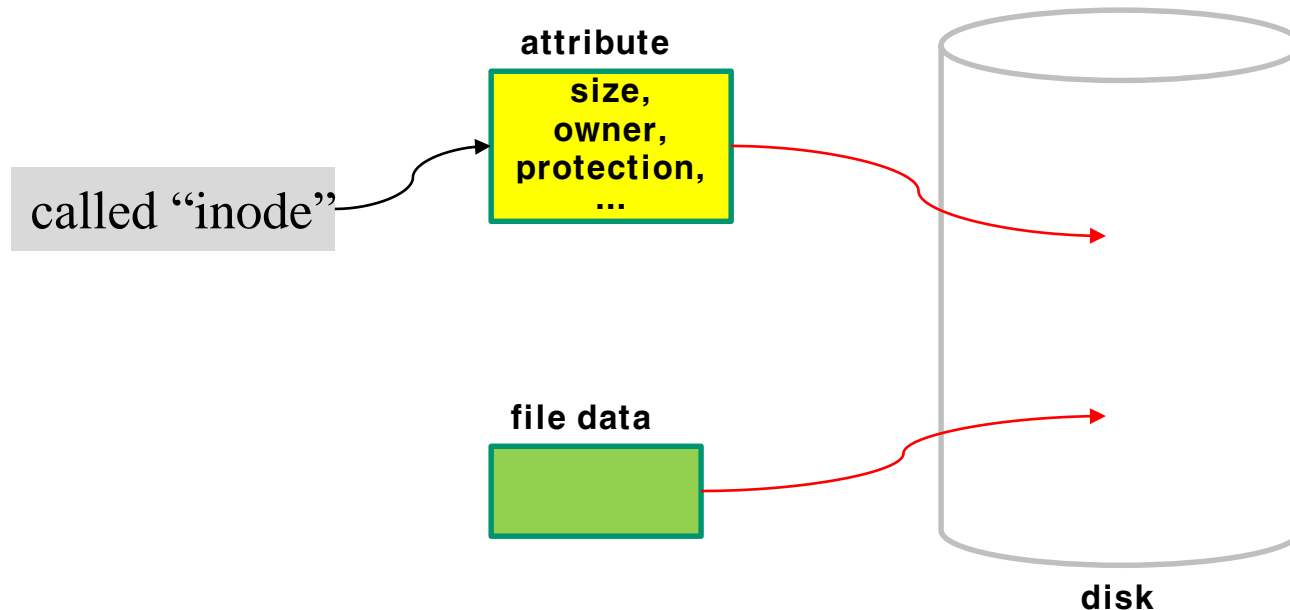
ls utility

-s: with -l, print size of each file, in blocks.

File system


File system

- Storing and retrieving “file data” & “file’s attribute”



File system


Mapping 3D of disk to 1D



File system

File system

- inode per each file



File system

Check inode

```
$ ls -il *
```

```
14951814 -rw-rw-r-- 1 kwon kwon 284 Sep  5 12:46 add.c  
14945425 -rwxrwxr-x 1 kwon kwon 8720 Sep  5 12:46 a.out  
14951666 -rw-rw-r-- 1 kwon kwon  81 Sep  5 12:23 hello.c  
14951659 -rw-rw-r-- 1 kwon kwon   0 Sep  5 12:06 test
```

```
$ stat add.c
```

```
File: 'add.c'
```

```
Size: 284          Blocks: 8          IO Block: 4096  regular file
```

```
Device: fc00h/64512d  Inode: 14951814  Links: 1
```

```
Access: (0664/-rw-rw-r--)  Uid: ( 1000/   kwon)  Gid: ( 1000/   kwon)
```

```
Access: 2018-09-05 12:46:33.198992197 +0900
```

```
Modify: 2018-09-05 12:46:30.515028545 +0900
```

```
Change: 2018-09-05 12:46:30.547028112 +0900
```


```
Birth: -
```

```
$
```

File system

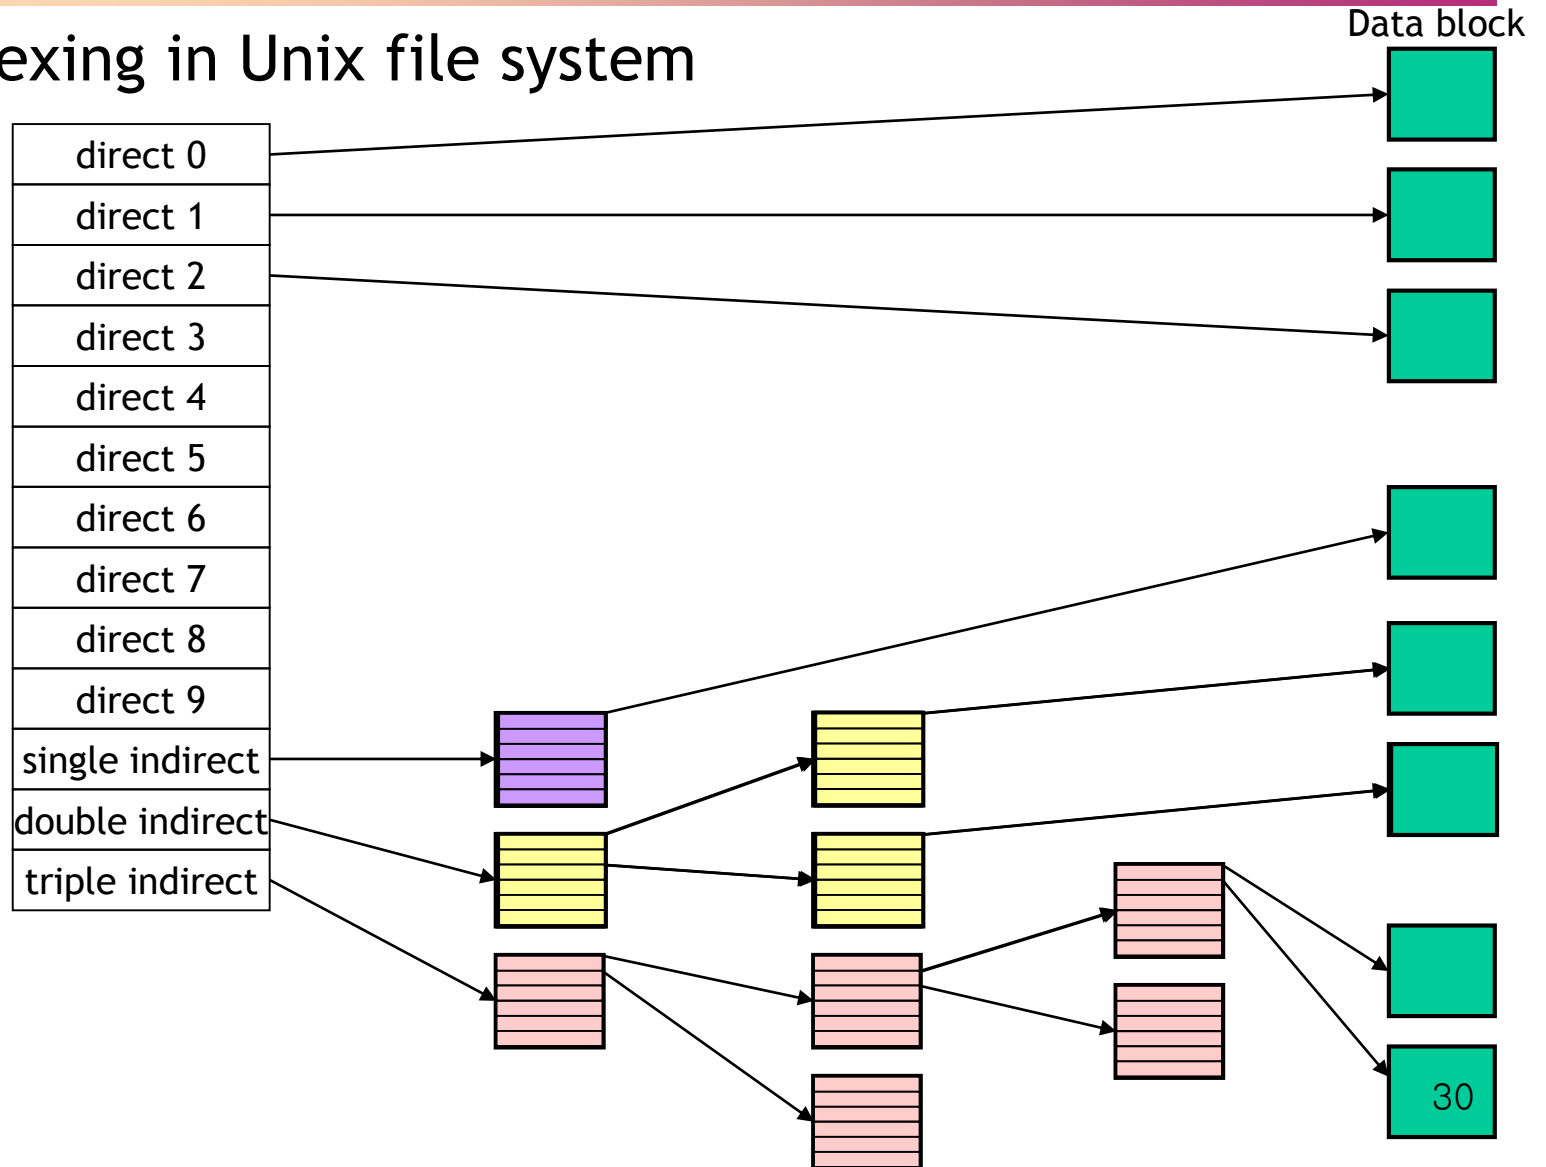
Example of Unix file system

- Separated : inode & data block
- Size of inode is constant
- ➔ Accessible by i-number(In the below, 0, 4, 6, and so on)



File system

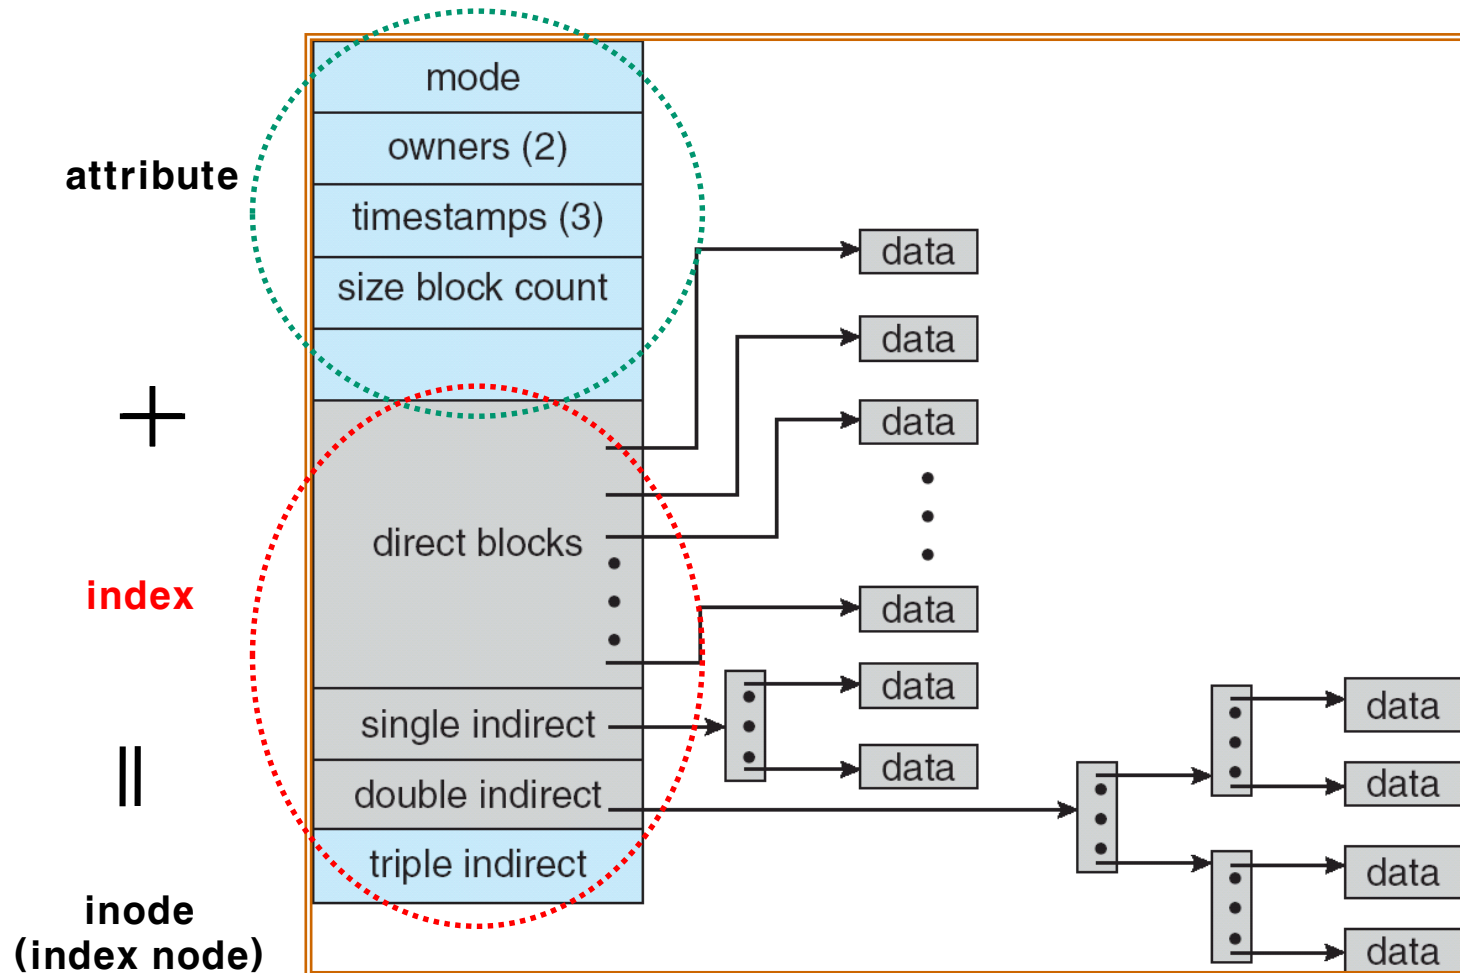
indexing in Unix file system



File system


inode in Unix file system

- Include attribute and a pointer to data block.



File system

Point to data block in inode



- Then, how to find inode?
 - Stored file name & inode number in Directory file.

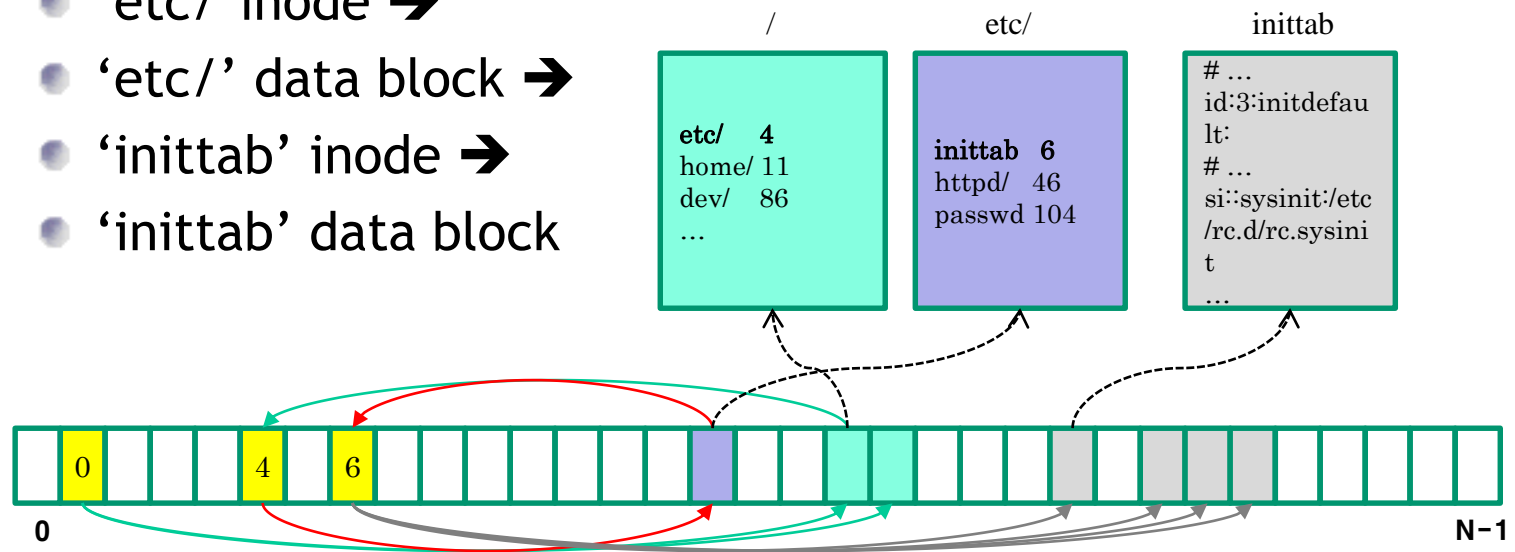
“/” directory

etc/	4
home/	11
dev/	86
...	

File system

Access to “/etc/inittab”?

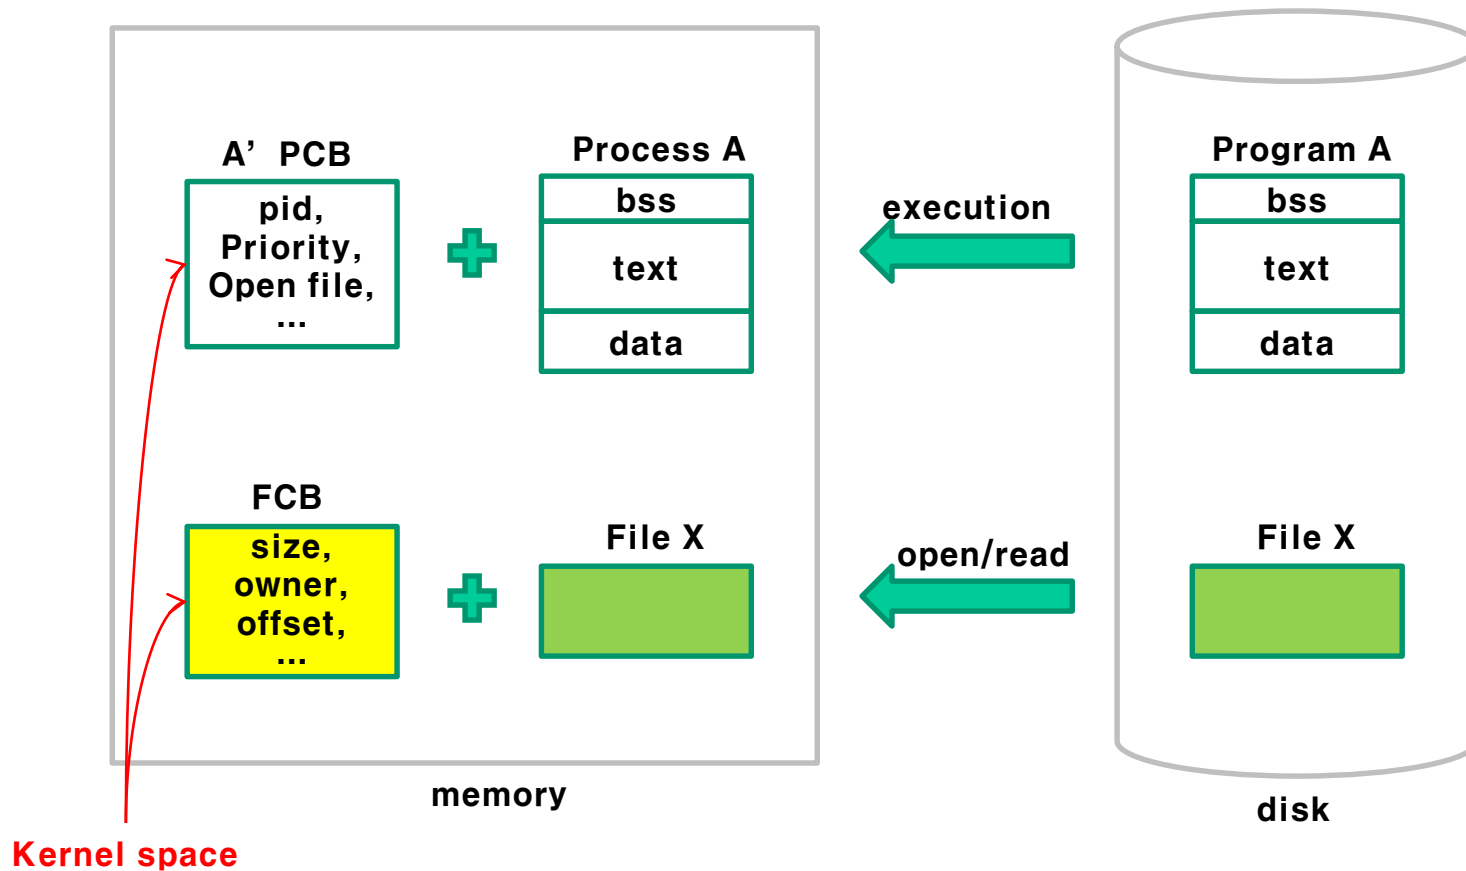
- ‘/’ inode →
- ‘/’ data block →
- ‘etc/’ inode →
- ‘etc/’ data block →
- ‘inittab’ inode →
- ‘inittab’ data block




- How to find ‘/’ inode?
 - Generally, i-number of ‘/’ is 0.

File in Process

- Kernel should manage metadata of file



File in Process

 FCB (File Control Block) : metadata to manage a file in kernel space

- size (e.g. 16KB)
- type (e.g. regular file)
- owner (e.g. obama)
- protection (e.g. rwxr--r--)
- Index to data block (e.g. sector address)
- device (e.g. /dev/hda0)
- Access location (e.g. **offset**)
- ...

File I/O system call review


 `fd = open("/etc/inittab", O_RDONLY);`



 `nread = read(3, buffer, BUF_SIZE);`

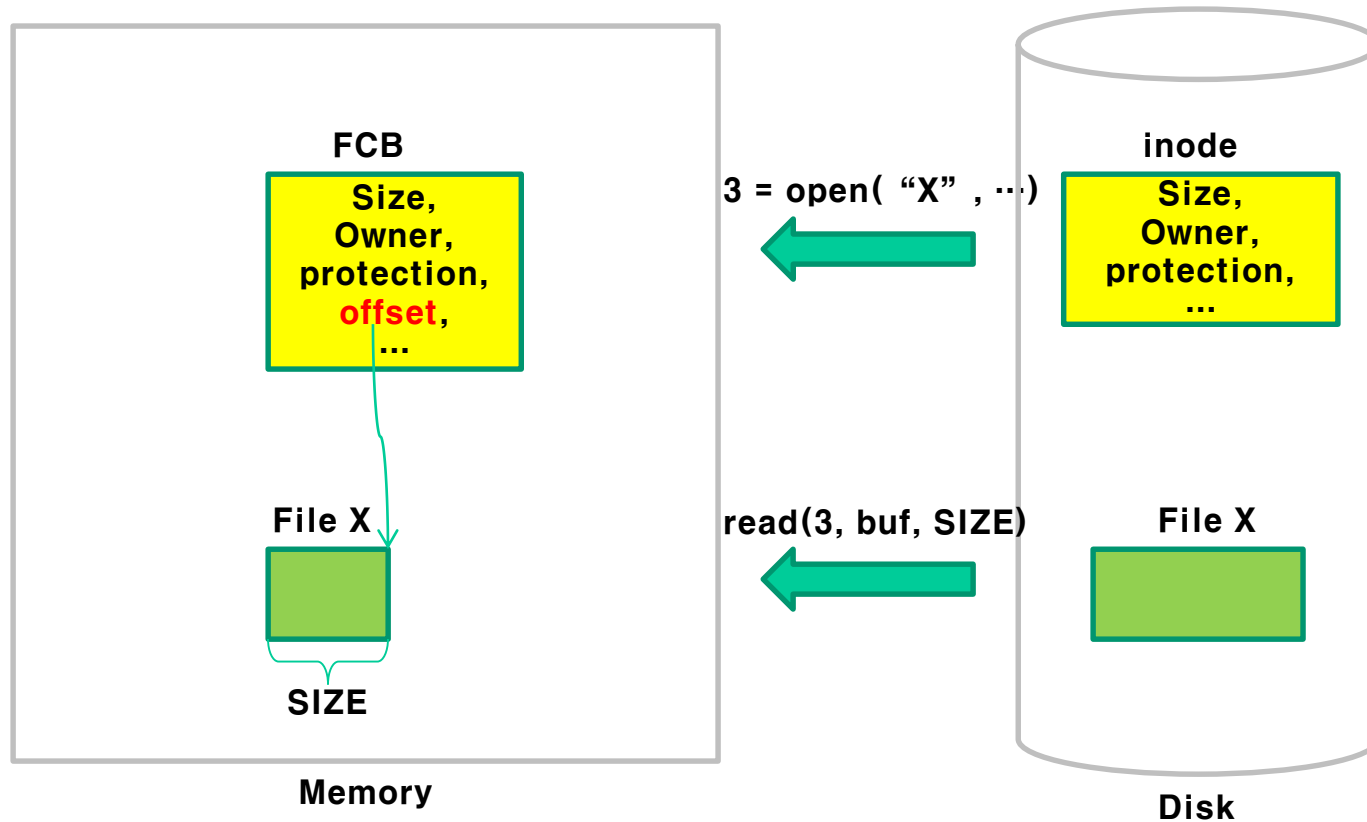


 `npos = lseek(3, 4096, SEEK_CUR);`



File I/O system call review

Open & Read

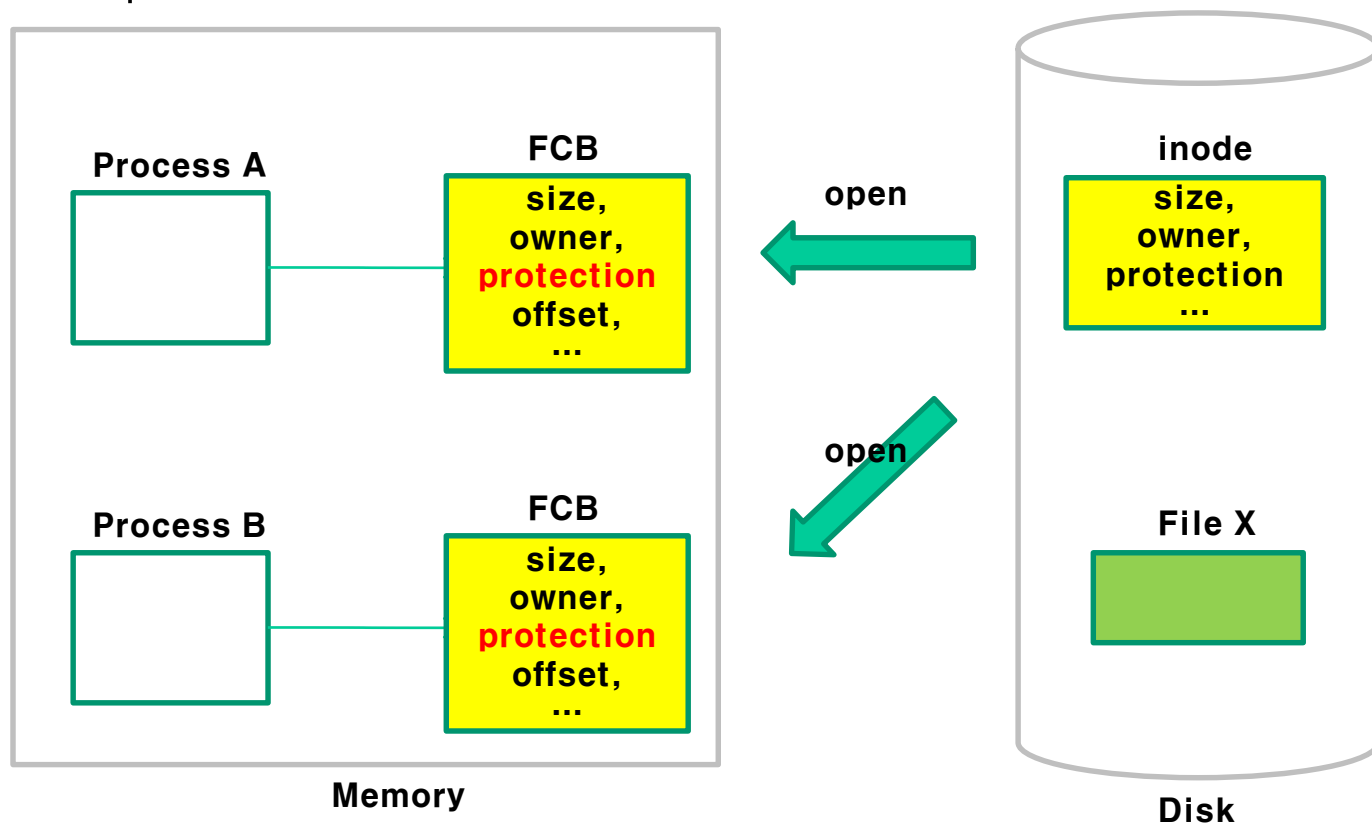


File in Process

❏ When two processes use a same file, two FCBs are needed

- In case that process A modifies 'permission'?

- \$ chmod a+w X



File in Process

What about modifying metadata and copying it?

- Inconsistency problem
- inefficient

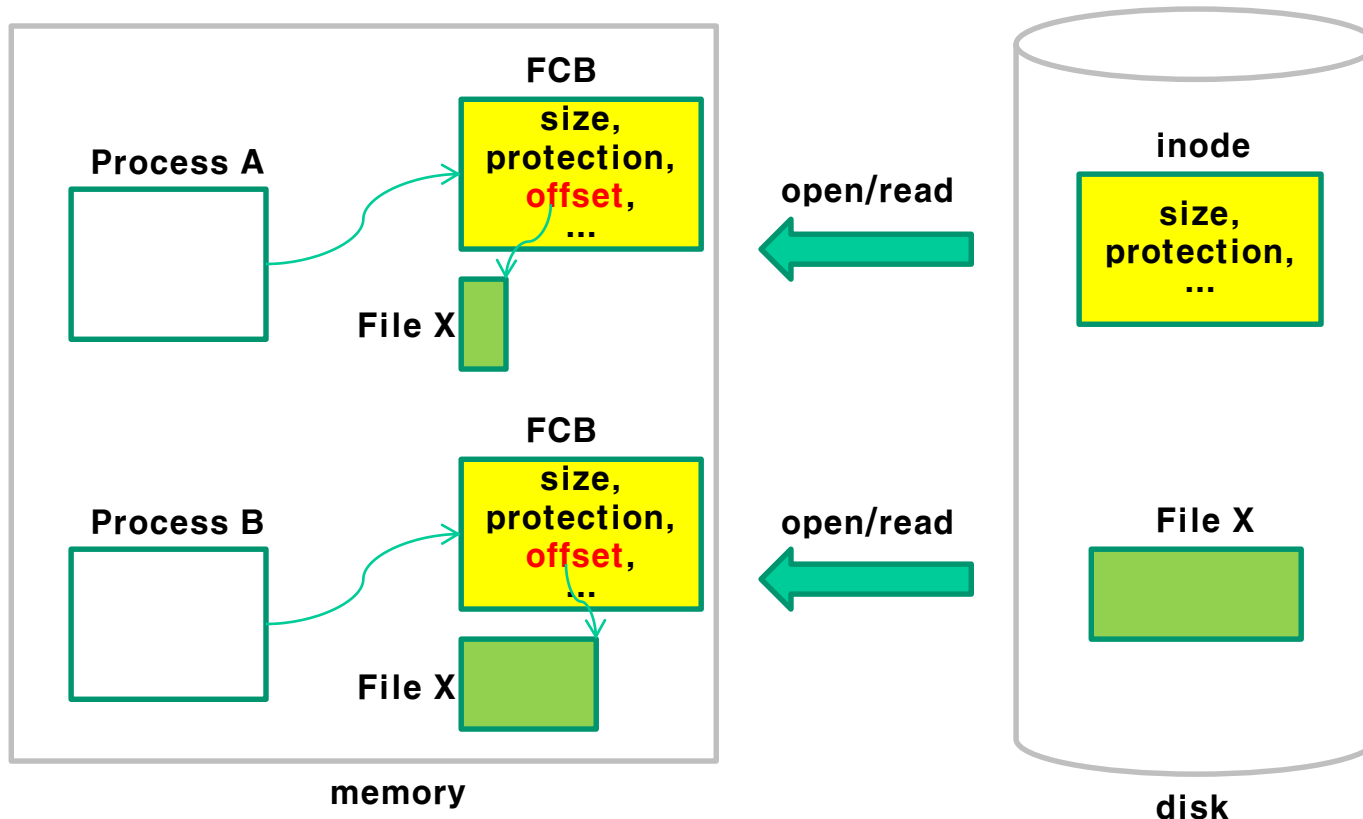
Share metadata

- Sharing permission, size, type, and so on, between processes
- What about offset?
 - Every process is reading/writing with different offset
 - Thus, individual process should have it

File in Process

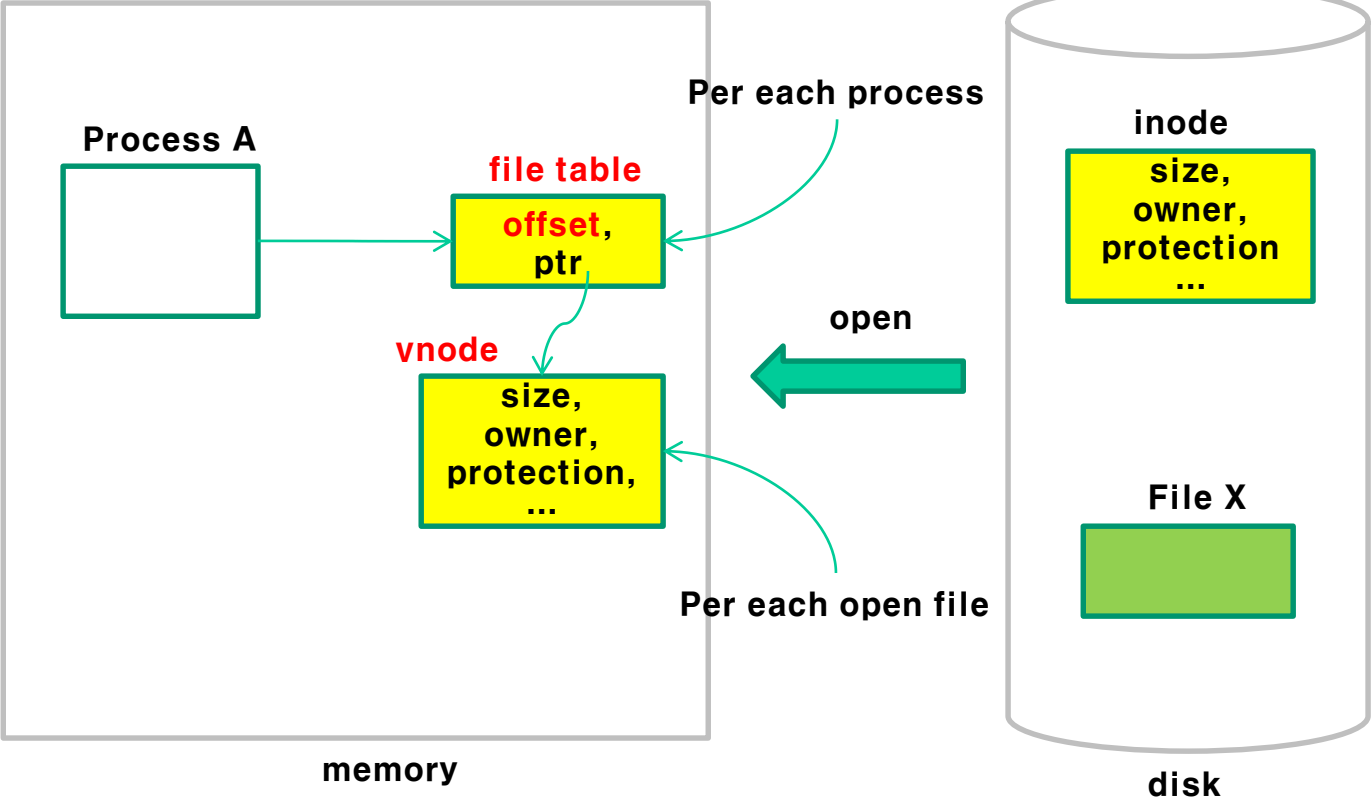
❏ Necessity of data structure division

- Share : protection, owner, ...
- Not-share : offset



File in Process

- Offset is stored in file table
- Others are stored in vnode



File in Process

file table

- Created whenever a file 'Open's
- Contents
 - offset
 - Pointer to vnode


File in Process

vnode

- Others except offset
- From inode in disk
 - protection mode
 - owner
 - size
 - time
 - data block location in disk

File in Process

- file descriptor table added
 - For a easy access after calling open()



File in Process

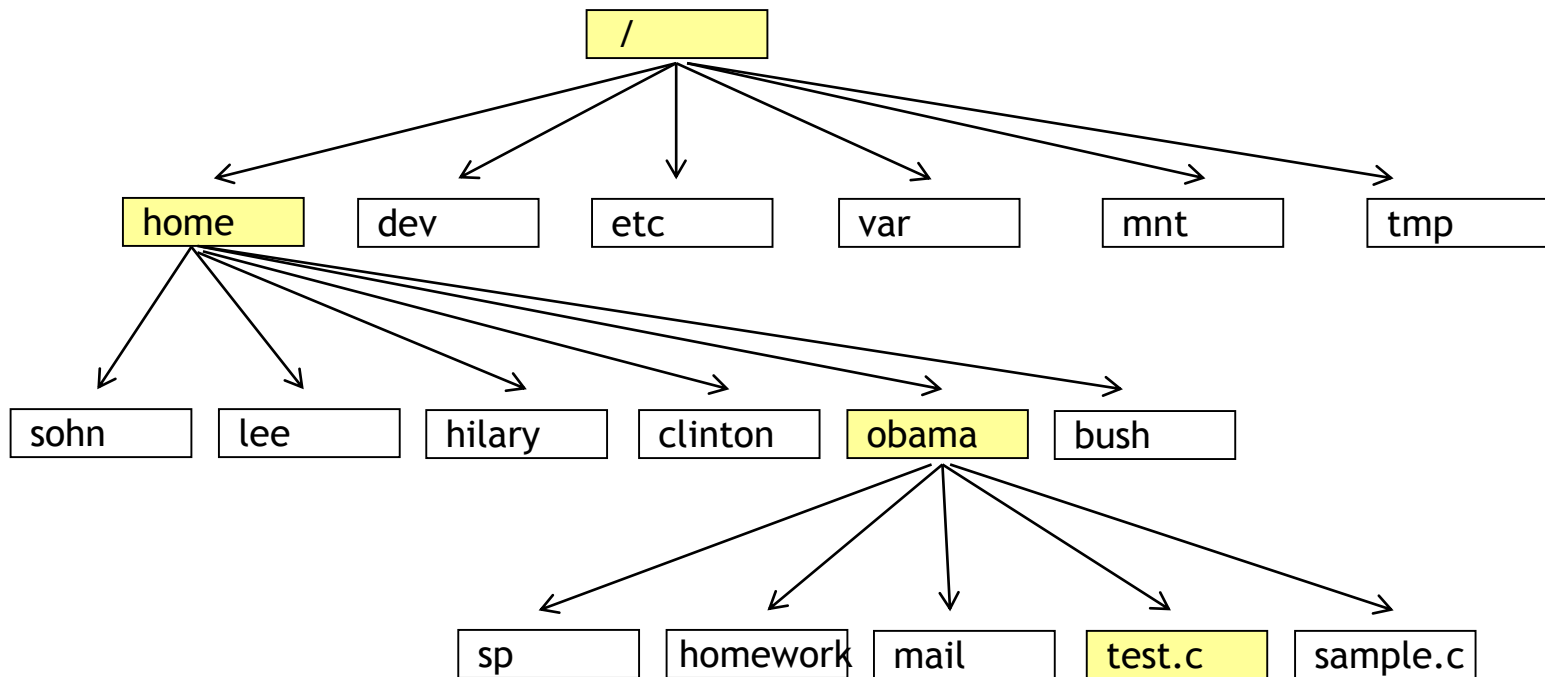
file descriptor table

- Per process data structure
- `fd = open("/a/b", ...)`
- `fd` is an index to access the file
 - Non-negative integers
 - 0, 1, 2 - standard input/output/error

open()

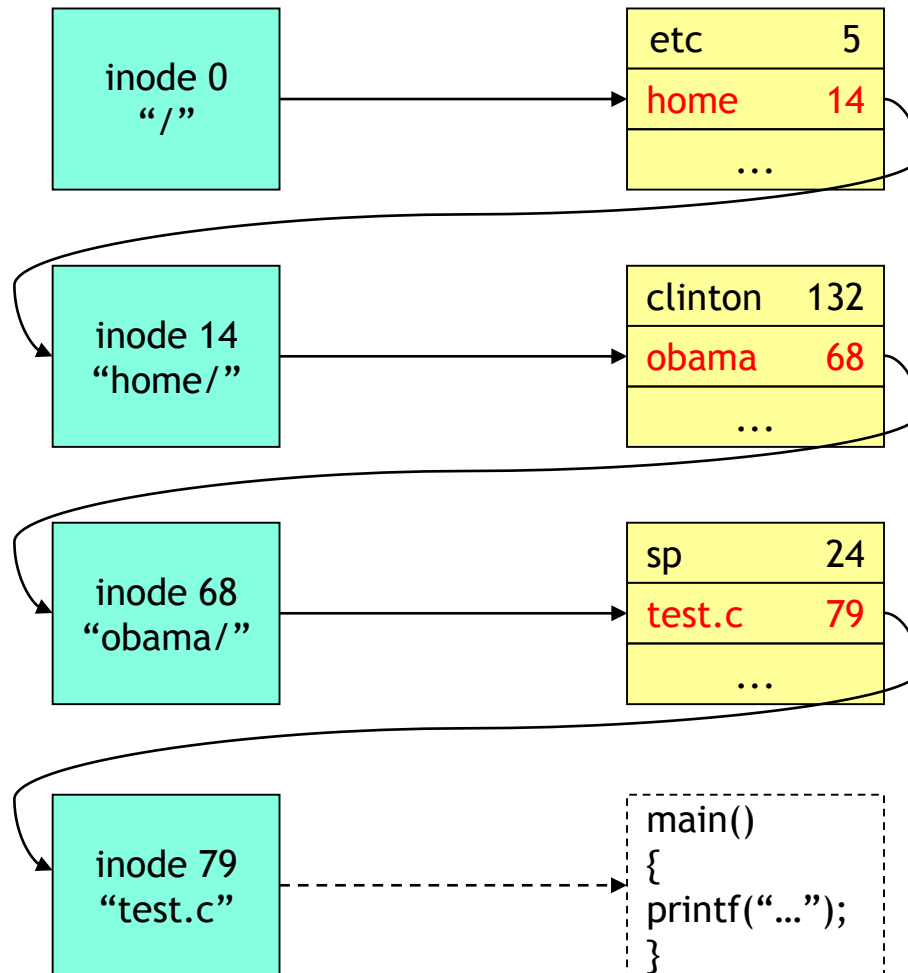
❏ \$ cat /home/obama/test.c

- open("/home/obama/test.c", O_RDONLY)




open()

1. Pathname lookup



open()

2. data structure creation for “test.c” in kernel
- Create vnode
 - Create file table (set offset to 0)
 - Create an entry in file descriptor table, and return file descriptor



open()

Pathname lookup overhead

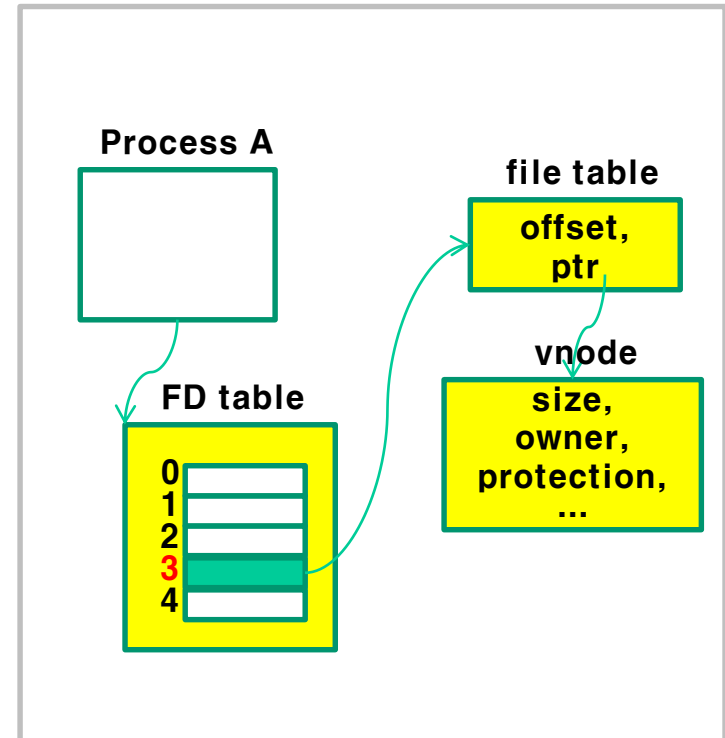
- open("/a/b", ...) needs many I/O operations
- Pathname lookup : executed one time!
- (pathname → file descriptor)
 - `fd = open("/a/b", ...)`
- Thereafter, system call uses a file descriptor instead of path
 - `read(fd, ...)`, `write(fd, ...)`, ...

File sharing

Three data structures in kernel when a process uses a file

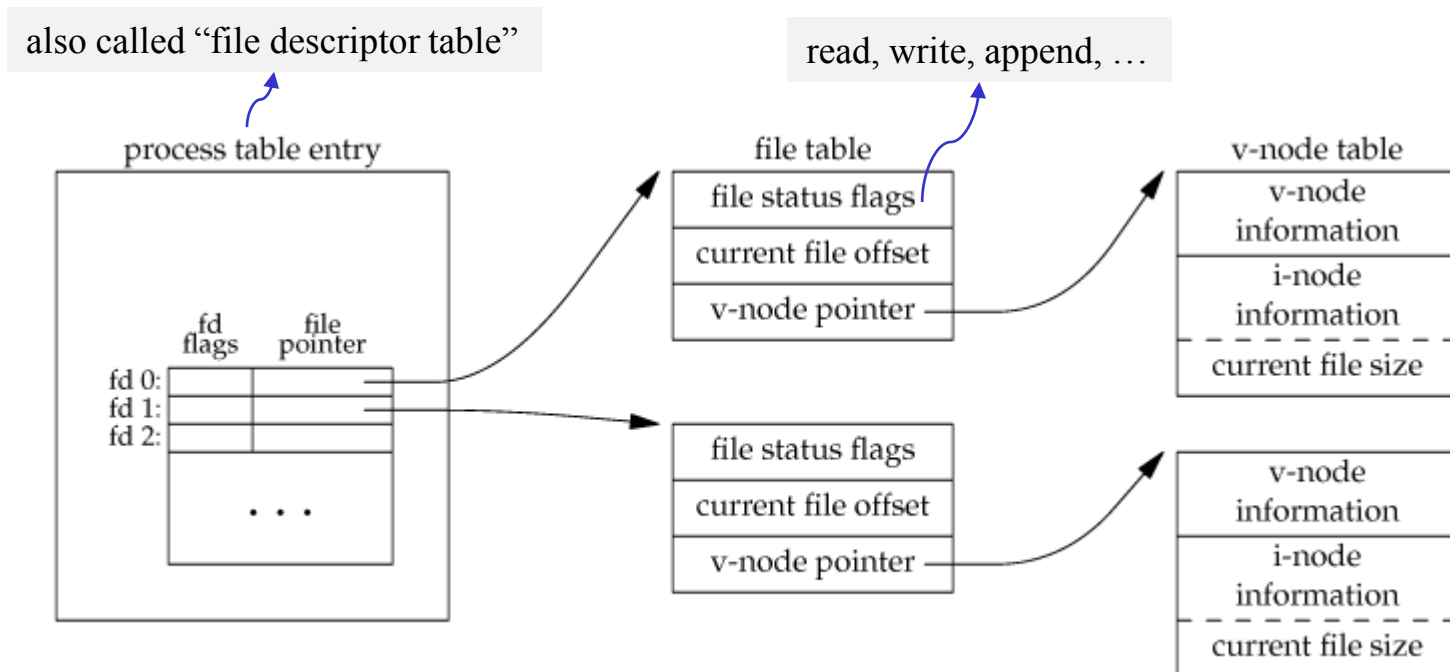
- File descriptor table
- File table
- vnode table (in linux, a generic-inode structure is used)

→ file sharing is possible



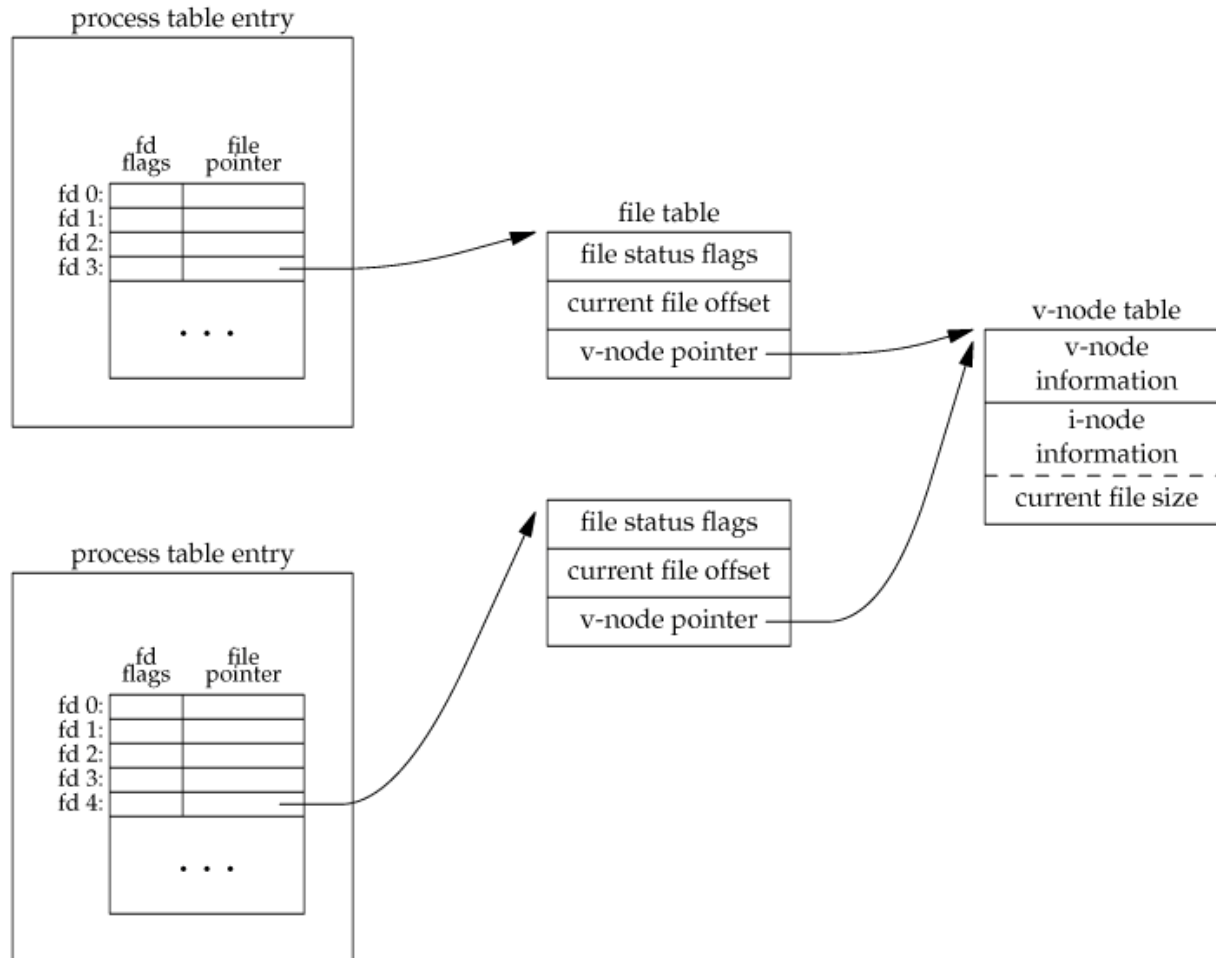
File sharing

- kernel data structures for a single process that has two different files open.



File sharing

Two independent processes with the same file open



dup() and dup2()

```
#include <unistd.h>
```

```
int dup(int fildes);
```

```
int dup2(int fildes, int fildes2);
```

Both return: new file descriptor if OK, -1 on error

dup

- create a copy of *fildes* and returns a new file descriptor specified by the lowest number available.

dup2

- makes *fildes2* be the copy of *fildes*, closing *fildes2* first if necessary.


Return values

- dup: the lowest numbered available file descriptor
- dup2: the new file descriptor with the *fildes2* argument

dup() and dup2()

Kernel data structures after “dup(1)”

- The next available descriptor is 3.




dup() and dup2()

Example

```
#include <unistd.h>
#include <fcntl.h>

int main(void)
{
    int fd;
    fd = creat("dup_result", 0644);
    dup2(fd, STDOUT_FILENO);
    close(fd);
    printf("hello world\n");
    return 0;
}
```



Execution

```
$ cat dup_result
hello world
```

sync(), fsync(), and fdatasync()

Delayed write

- When write data to a file, the data is copied into buffers.
- The data is physically written to disk at some later time.

When the delayed-write blocks are written to disk?

- Buffer is filled with the delayed-write blocks or
- Periodically by update daemon (usually every 30 seconds)

sync(), fsync(), and fdatasync()

```
#include <unistd.h>
```

```
int fsync(int fildes);
```

```
int fdatasync(int fildes);
```

Returns: 0 if OK, -1 on error

```
void sync(void);
```

sync

- Write **all** the modified buffer blocks to disk.

fsync

- Write only the modified (**data + attribute**) buffer blocks of a single file.

fdatasync

- Write only the modified **data** buffer blocks of a single file.

fcntl()

```
#include <fcntl.h>
```

```
int fcntl(int filedes, int cmd, ... /* int arg */);
```

Returns: depends on cmd if OK (see following), -1 on error

 Change the properties of a file that is already open

- Duplicate an existing descriptor (cmd = F_DUPFD)
- Get/set file descriptor flags (cmd = F_GETFD or F_SETFD)
- Get/set file status flags (cmd = F_GETFL or F_SETFL)
- Get/set asynchronous I/O ownership (cmd = F_GETOWN or F_SETOWN)
- Get/set record locks (cmd = F_GETLK, F_SETLK, or F_SETLKW)

fcntl()

example

```
/* omitted header files */

int main()
{
    int mode, fd, value;

    fd = open("test.sh", O_RDONLY|O_CREAT);
    value = fcntl(fd, F_GETFL, 0);

    mode = value & O_ACCMODE;
    if (mode == O_RDONLY)
        printf("O_RDONLY setting\n");
    else if (mode == O_WRONLY)
        printf("O_WRONLY setting\n");
    else if (mode == O_RDWR)
        printf("O_RDWR setting\n");
}
```

fcntl()

Execution

```
$ ./fgetfl_test  
O_RDONLY setting  
$
```