



Process control



Process

Process

- a program in execution
- program image + environment
 - `text(code)` / `data` / `stack` / `heap` segment
 - kernel data structure, address space, open files, ...

Process

```
/* test_program.c */
#include <stdio.h>
int a,b;
int glob_var = 3;

void test_func(void) {
    int local_var, *buf;

    buf = (int *) malloc(10,000 * sizeof(int));
    ...
}

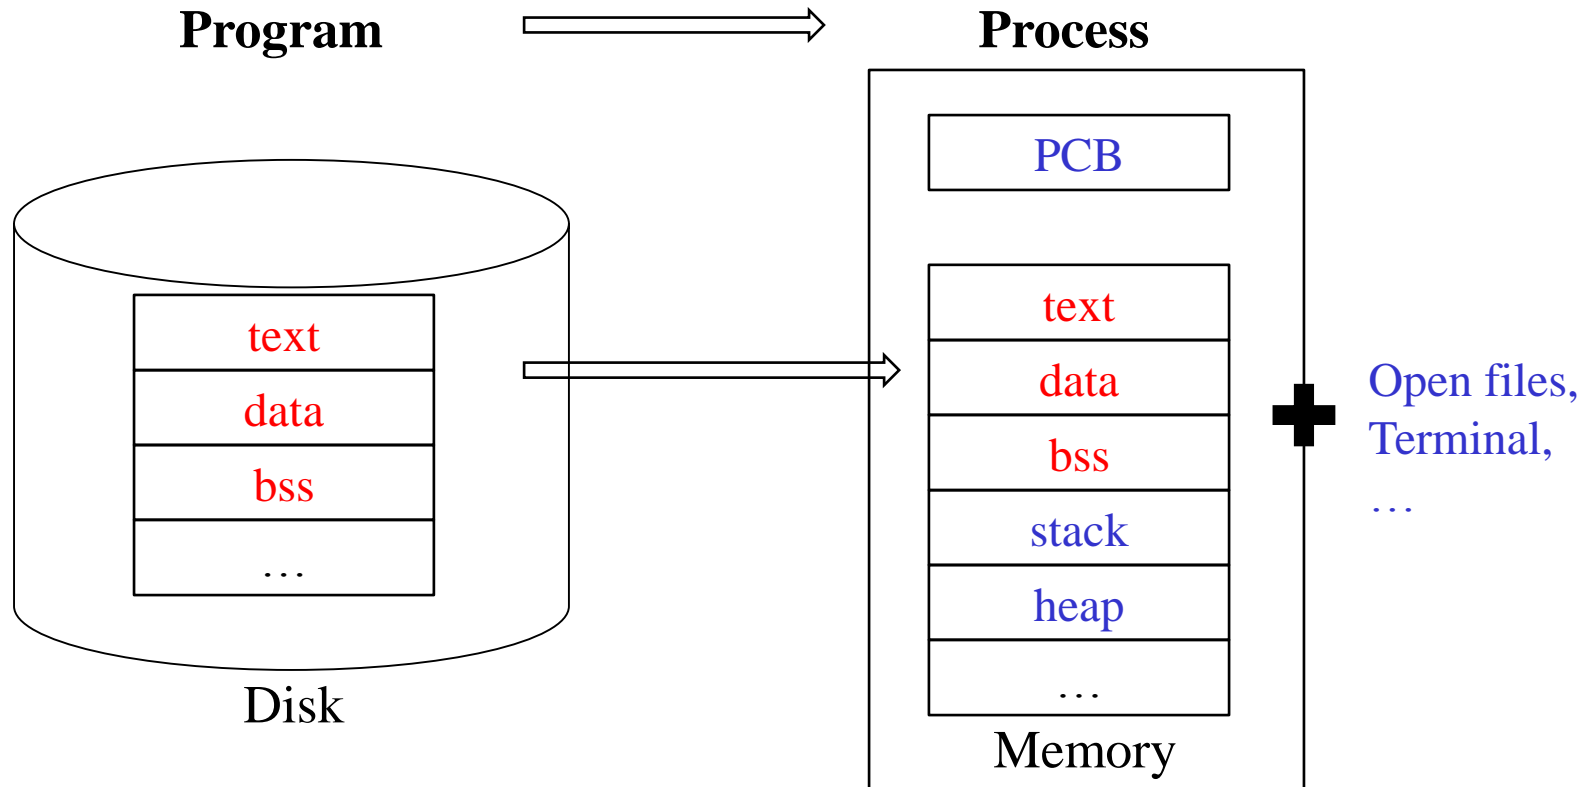
int main(int argc, char *argv[]) {
    int i = 1;
    int local_var;

    a = i + 1;
    printf("value of a = %d\n", a);
    test_func();
}
```

```
$ gcc -o test_program test_program.c
```

Process

 \$./test_program



* bss: Block Started by Symbol
PCB: Process Control Block

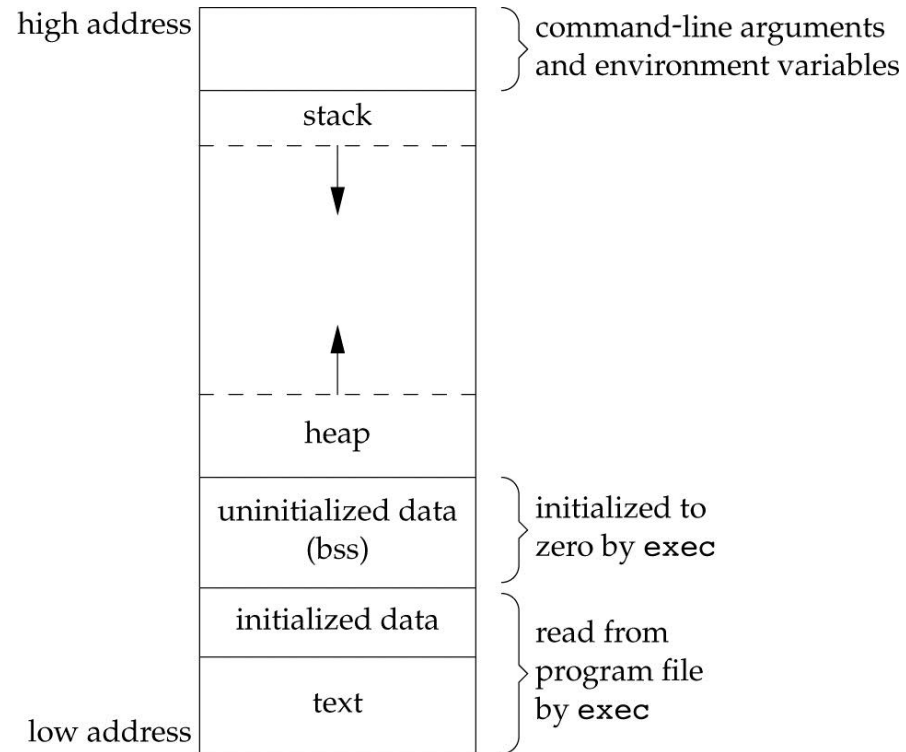


Figure 7.6 Typical memory arrangement

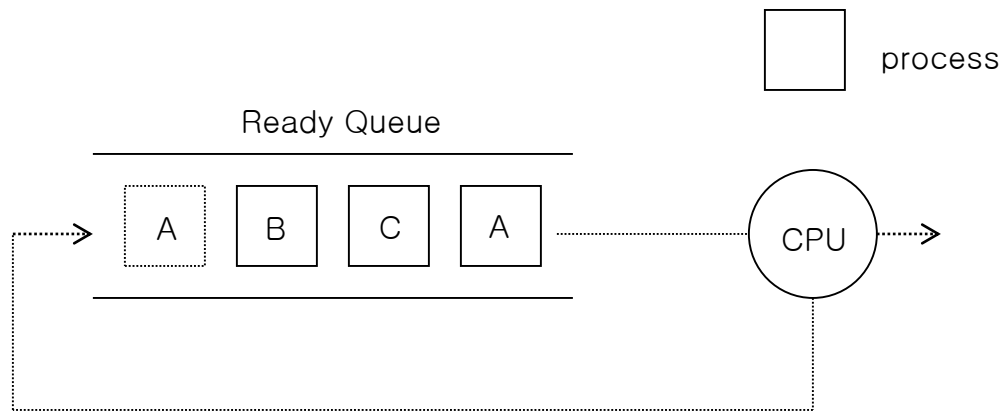
Process

 UNIX is a multitasking system.

```
$ ps -e
PID TTY      TIME CMD
1 ?        00:00:00 init
2 ?        00:00:00 migration/0
3 ?        00:00:00 ksoftirqd/0
4 ?        00:00:00 watchdog/0
5 ?        00:00:00 migration/1
6 ?        00:00:00 ksoftirqd/1
7 ?        00:00:00 watchdog/1
8 ?        00:00:00 migration/2
...
29122 pts/9   00:00:00 bash
29151 ?        00:00:00 sshd
29170 ?        00:00:00 sshd
29171 pts/10    00:00:00 bash
30465 ?        00:00:00 httpd
$
```

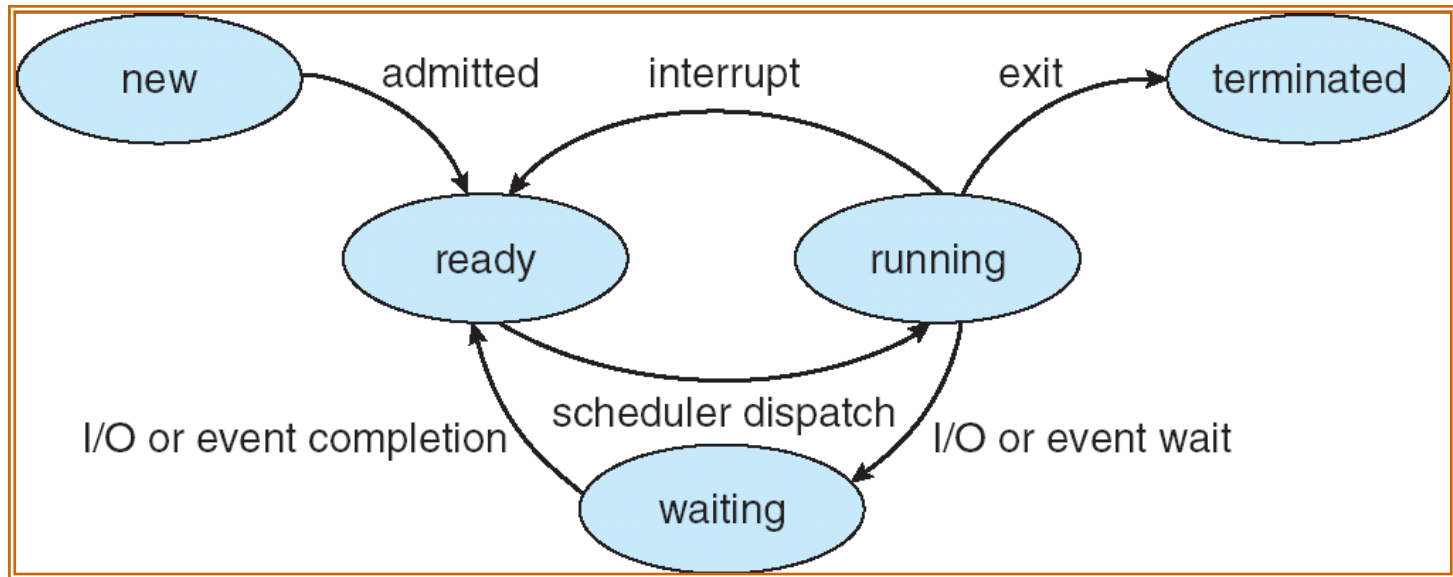
Process

🖥️ UNIX is a time sharing system.



Process

State of a process



Process identifiers

process id

- Every process has a unique process ID.
- As processes terminate, IDs become candidates for reuse.

```
$ ps -e
PID TTY          TIME CMD
1 ?           00:00:00 init
2 ?           00:00:00 migration/0
3 ?           00:00:00 ksoftirqd/0
4 ?           00:00:00 watchdog/0
5 ?           00:00:00 migration/1
...
29122 pts/9    00:00:00 bash
29151 ?         00:00:00 sshd
29170 ?         00:00:00 sshd
29171 pts/10    00:00:00 bash
30465 ?         00:00:00 httpd
$
```

Process identifiers

Process 0 (scheduler process or swapper)

- Part of the kernel.
- System process.

Process 1 (init process)

- /sbin/init
- Invoked at the end of the bootstrap procedure.
- Initialize the UNIX system with /etc/rc* or /etc/inittab.
- Never dies.
- A user process with superuser privilege.

Process identifiers

```
#include <unistd.h>
```

```
pid_t getpid(void);
```

Returns: process ID of calling process

```
pid_t getppid(void);
```

Returns: parent process ID of calling process

```
uid_t getuid(void);
```

Returns: real user ID of calling process

```
uid_t geteuid(void);
```

Returns: effective user ID of calling process

```
gid_t getgid(void);
```

Returns: real group ID of calling process

```
gid_t getegid(void);
```

Returns: effective group ID of calling process

fork()

```
#include <unistd.h>
```

```
pid_t fork(void);
```

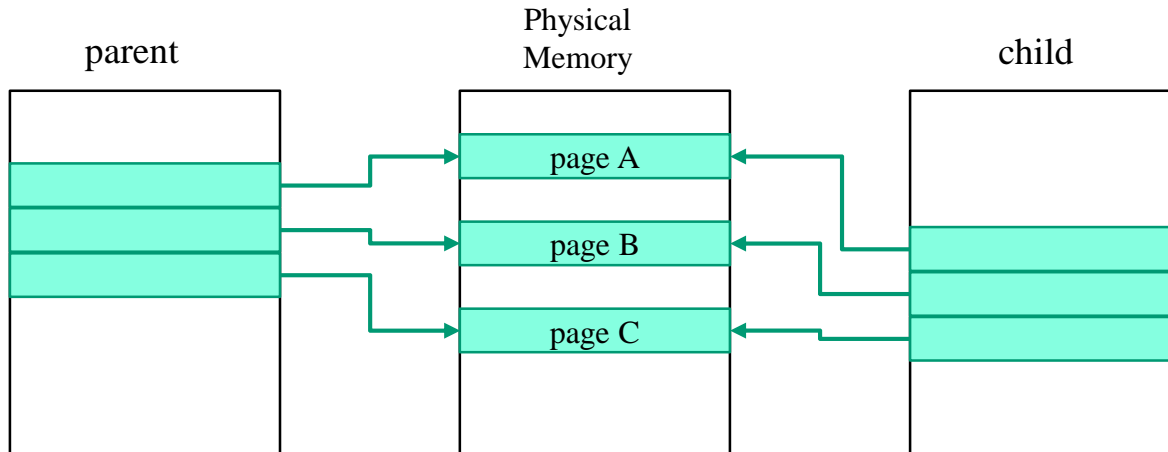
Returns: 0 in child, process ID of child in parent, -1 on error

Create a child process.

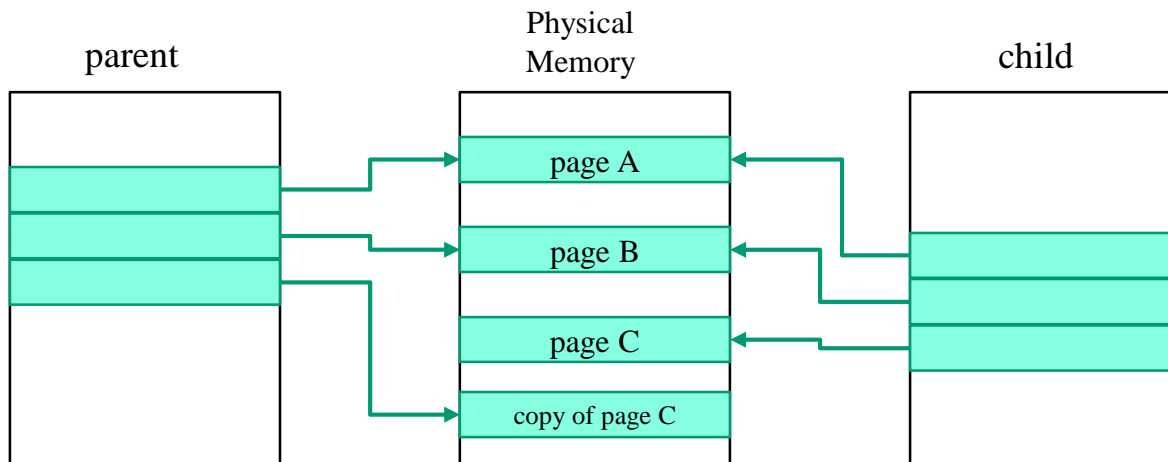
- It is called once but returns twice.
 - Return value of child: 0
 - Return value of parent: PID of child
- Both child and parent continue executing with the instruction that follows the call to fork().
- The child is a copy of parent.
 - Child gets a copy of the parent's data, heap, and stack.
 - Parent and child often share the text segment.

fork()

Copy on write



- After parent process write a page C.



fork()

Example

```
#include "apue.h"

int  glob = 6;    /* external variable in initialized data */
char buf[] = "a write to stdout\n";

int
main(void)
{
    int    var;    /* automatic variable on the stack */
    pid_t  pid;

    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n");    /* we don't flush stdout */
```

What about strlen() instead of sizeof()?

fork()

Example(cont.)

```
if ((pid = fork()) < 0) {
    err_sys("fork error");
} else if (pid == 0) {           /* child */
    glob++;                     /* modify variables */
    var++;
} else {
    sleep(2);                   /* parent */
}

printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
exit(0);
}
```

fork()

```
#include "apue.h"
int  glob = 6;
char buf[] = "a write to stdout\n";

int main(void)
{
    int  var;
    pid_t pid;

    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n"); /* we don't flush stdout */

    if ((pid = fork()) < 0) { /* pid is child's pid(non-0). */
        err_sys("fork error");
    } else if (pid == 0) { /* child */
        glob++; /* modify variables */
        var++;
    } else { /* parent */
        sleep(2);
    }

    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit(0);
}
```

Parent process

```
#include "apue.h"
int  glob = 6;
char buf[] = "a write to stdout\n";

int main(void)
{
    int  var;
    pid_t pid;

    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n"); /* we don't flush stdout */

    if ((pid = fork()) < 0) { /* pid is 0. */
        err_sys("fork error");
    } else if (pid == 0) { /* child */
        glob++; /* modify variables */
        var++;
    } else { /* parent */
        sleep(2);
    }

    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit(0);
}
```

Child process

fork()

execution

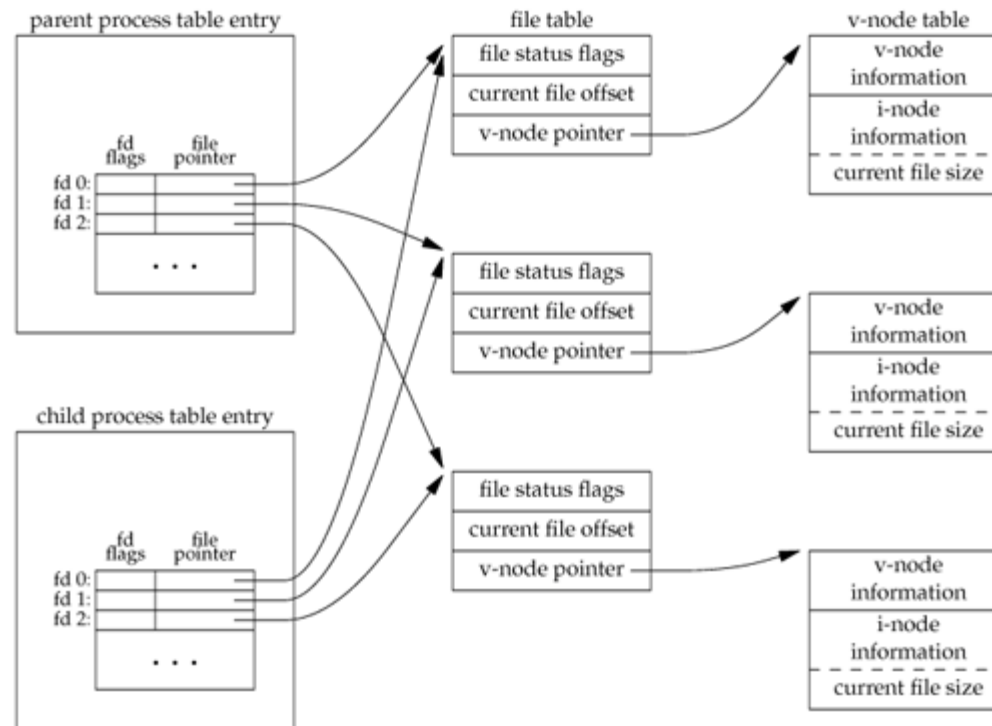
```
$ ./a.out
a write to stdout
before fork
pid = 430, glob = 7, var = 89
pid = 429, glob = 6, var = 88
$ ./a.out > temp.out
$ cat temp.out
a write to stdout
before fork
pid = 432, glob = 7, var = 89
before fork
pid = 431, glob = 6, var = 88
$
```

We never know if the child starts executing before the parent or vice versa.
- It depends on the process scheduling algorithm.

fork()

File sharing

- Sharing of open files between parent and child after fork()



fork()

Information shared by child and parent.

- real-uid(gid), effective-uid(gid)
- controlling terminal
- current working directory, root directory
- signal handlers
- environment
- resource limits

differences between child and parent.

- the return value from fork.
- PID and PPID
- child's resource utilizations are set to 0.
- pending signals

fork()

Two main reasons for fork to fail

- if there are already too many processes in the system.
- if the total number of processes for this real user ID exceeds the system's limit.

Two uses for fork

- A process wants to duplicate itself.
 - Parent and child can each execute different sections of code at the same time.
 - Common for **network servers**.
- A process wants to execute a different program.
 - Common for shells.
 - Child does **exec()** right after it returns from **fork()**.

vfork()

```
#include <unistd.h>
```

```
pid_t vfork(void);
```

Returns: 0 in child, process ID of child in parent, -1 on error

- same calling sequence and same return values as fork().
- intended to create a new process when the purpose of the new process is to exec() a new program.
 - Does not copy the address space of parent into the child.
 - The child calls exec() or exit() right after the vfork().
 - The child runs in the address space of the parent.
 - Provides an efficiency.
- vfork() guarantees that the child runs first.

vfork()

Example

```
#include "apue.h"
int  glob = 6;    /* external variable in initialized data */

int
main(void)
{
    int  var;      /* automatic variable on the stack */
    pid_t pid;

    var = 88;
    printf("before vfork\n"); /* we don't flush stdio */
    if ((pid = vfork()) < 0) {
        err_sys("vfork error");
    } else if (pid == 0) { /* child */
        glob++;           /* modify parent's variables */
        var++;
        _exit(0);         /* child terminates */
    }
}
```

vfork()

Example(cont.)

```
/*  
 * Parent continues here.  
 */  
printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);  
exit(0);  
}
```

Execution

```
$ ./a.out  
before vfork  
pid = 29039, glob = 7, var = 89  
$
```

exit()

```
#include <stdlib.h>
void exit(int status);
void _Exit(int status);

#include <unistd.h>
void _exit(int status);
```

❏ causes **normal program termination** and the value of status is returned to the parent.

❏ **_exit and _Exit**

- return to the kernel immediately.
- **_exit** : System call

❏ **exit**

- C standard library
- performs cleanup processing and then returns to the kernel.
 - all open streams are flushed and closed.

atexit()

```
#include <stdlib.h>
int atexit(void (*func)(void));
```

Returns: 0 if OK, nonzero on error

exit handler

- Function automatically called by exit
- Registered by atexit
- exit() calls the registered handlers in reverse order of their registration with repetition.

Process termination

Two types of process termination

- (voluntarily) **normal** termination
 - return from main (equivalent to calling exit)
 - calling exit
 - calling _exit or _Exit
 - Returning of the last thread from its start routine
 - Calling pthread_exit from the last thread
- (involuntarily) **abnormal** termination
 - calling abort
 - Receipt of a signal
 - Response of the last thread to a cancellation request

Regardless of how a process terminate, the same code is eventually executed.

- Close open descriptors, release the memory, ...

Process termination

Termination status

- Terminating process notify its parent how it terminated.
- normal termination
 - pass an exit status as argument to `exit()` or `_exit()`.
 - return value from `main()`
 - exit status is converted into a termination status.
- abnormal termination
 - Kernel generates a termination status to indicate the reason for the abnormal termination.
- The parent of the terminated process can obtain the termination status from **`wait()`** or **`waitpid()`**.

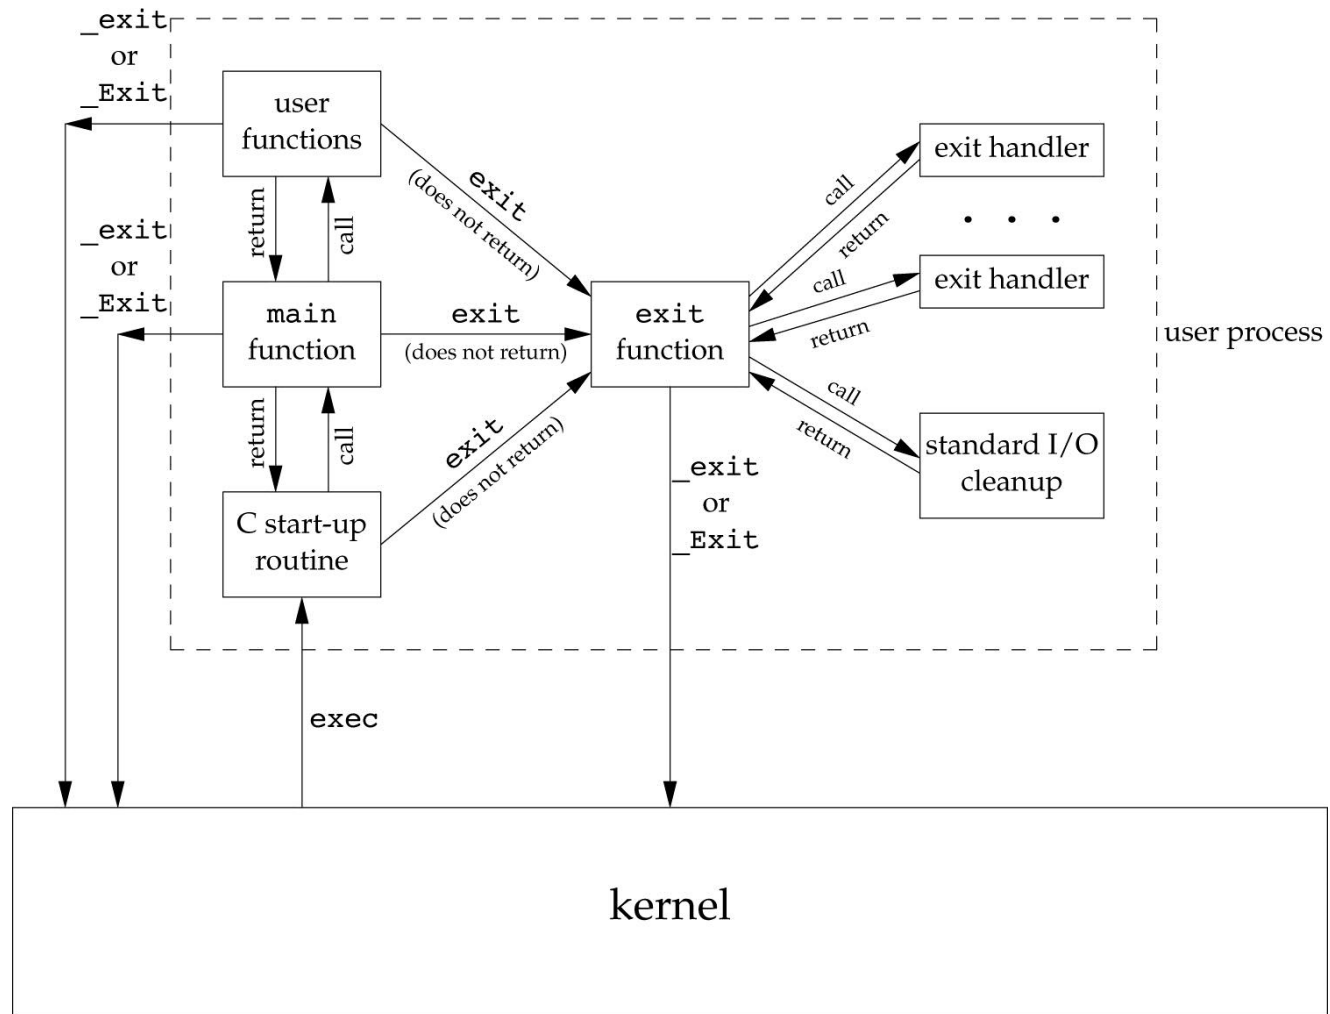


Figure 7.2 How a C program is started and how it terminates

Process termination

zombie state

- Kernel has to keep some information
 - PID, termination status, and CPU usage time
 - This information is available when parent calls wait().
- the process that has terminated, but whose parent has not yet waited for it, is called a **zombie**.

If the parent terminates before the child?

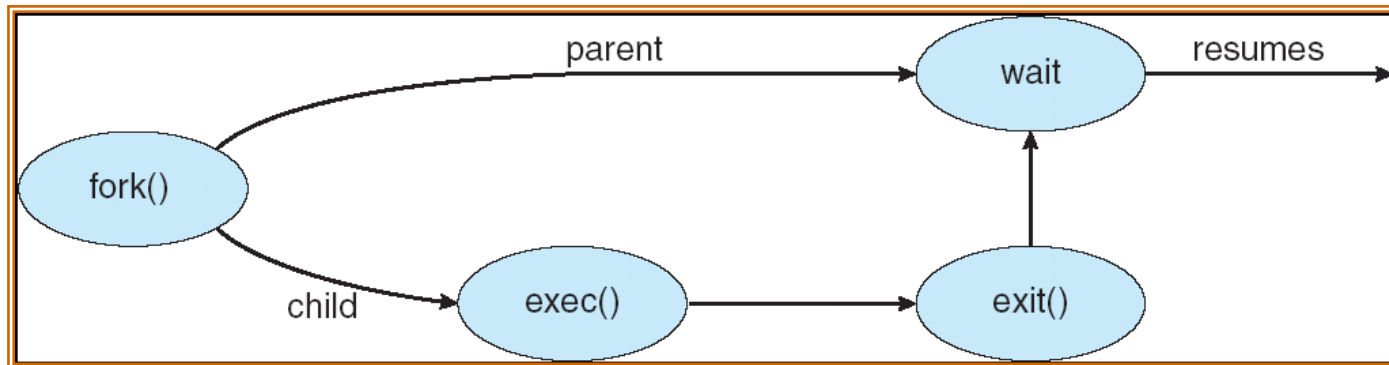
- **init** process becomes the parent of the orphaned process.

If an orphaned process is terminated?

- **init** calls **wait()** to fetch the termination status.

wait() and waitpid()

- ❏ A process that calls wait() or waitpid()
- Block, if all of its children are still running.
 - Return immediately if a child has terminated.
 - Return immediately with an error, if it doesn't have any child processes.



wait() and waitpid()

```
#include <sys/wait.h>
```

```
pid_t wait(int *statloc);
```

Both return: process ID if OK, 0 (see later), or -1 on error

wait()

- If a child has already terminated and is a zombie, wait returns immediately with that child's status.
- Otherwise, it blocks the caller until a child terminates.
- If the caller blocks and has multiple children, wait returns when one terminates.

statloc argument

- Store termination status in the location pointed to by *statloc*.

wait() and waitpid()

Macro to examine the termination status

- **WIFEXITED(status)**
 - True if a child terminated normally
- **WEXITEDSTATUS(status)** (only if WIFEXITED(status) is true)
 - Fetch the low-order 8 bits of the argument the child passed
- **WIFSIGNALED(status)**
 - True if a child terminated abnormally by receipt of a signal
- **WTERMSIG(status)** (only if WIFSIGNALED(status) is true)
 - Fetch the signal number
- **WIFSTOPPED(status)**
 - True if a child is currently stopped
- **WSTOPSIG(status)** (only if WIFSTOPPED(status) is true)
 - Fetch the signal number

wait() and waitpid()

Example

```
#include <sys/types.h>
#include <sys/wait.h>

void pr_exit(int status)
{
    if (WIFEXITED(status))
        printf("normal termination, exit status = %d\n",
               WEXITSTATUS(status));
    else if (WIFSIGNALED(status))
        printf("abnormal termination, signal number = %d\n",
               WTERMSIG(status));
    else if (WIFSTOPPED(status))
        printf("child stopped, signal number = %d\n",
               WSTOPSIG(status));
}
```

wait() and waitpid()

Example(cont.)

```
#include <sys/types.h>
#include <sys/wait.h>

int main(void)
{
    pid_t  pid;
    int    status;

    if ((pid = fork()) < 0)
        printf("fork error");
    else if (pid == 0)                /* child */
        exit(7);

    if (wait(&status) != pid)        /* wait for child */
        printf("wait error");
    pr_exit(status);                 /* and print its status */
}
```

wait() and waitpid()

Example(cont.)

```
if ( (pid = fork()) < 0)
    printf("fork error");
else if (pid == 0)           /* child */
    abort();                 /* generates SIGABRT */

if (wait(&status) != pid)    /* wait for child */
    printf("wait error");
pr_exit(status);             /* and print its status */

if ( (pid = fork()) < 0)
    printf("fork error");
else if (pid == 0)           /* child */
    status /= 0;             /* divide by 0 generates SIGFPE */

if (wait(&status) != pid)    /* wait for child */
    printf("wait error");
pr_exit(status);             /* and print its status */
}
```

wait() and waitpid()

Execution

```
$ ./a.out
```

```
normal termination, exit status = 7
```

```
abnormal termination, signal number = 6
```

```
abnormal termination, signal number = 8
```

```
$
```

Each signal number is defined in <signal.h>

SIGABRT: 6

SIGFPE: 8

wait() and waitpid()

```
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int *statloc, int options);
```

Both return: process ID if OK, 0 (see later), or -1 on error

Why waitpid() is required?

- If we have more than one child, **wait() returns on termination of any of the children.**
- What if we want to wait for a specific process to terminate?

waitpid()

- Waits for a specific process with *pid* argument.
- Provides some controls with *options* argument.

wait() and waitpid()

pid argument

- `pid == -1`
 - Waits for any child process.
 - In this case, `waitpid()` is equivalent to `wait()`.
- `pid > 0`
 - Waits for the child whose process ID equals `pid`.
- `pid == 0`
 - Waits for any child whose process group ID equals that of the calling process.
- `pid < -1`
 - Waits for any child whose process group ID equals the absolute value of `pid`.

wait() and waitpid()

options argument

- WNOHANG
 - Not block(return immediately) if a child specified by pid is not terminated.
- WUNTRACED
 - Not block if a child specified by pid has stopped.

wait() and waitpid()

Example

```
#include "apue.h"
#include <sys/wait.h>

int main(void)
{
    pid_t  pid;

    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid == 0) {          /* first child */
        if ((pid = fork()) < 0)
            err_sys("fork error");
        else if (pid > 0)
            exit(0);               /* parent from second fork == first child */
        /* We're the second child; our parent becomes init. */
        sleep(2);
        printf("second child, parent pid = %d\n", getppid());
        exit(0);
    }
}
```


wait() and waitpid()

Example

```
if (waitpid(pid, NULL, 0) != pid) /* wait for first child */
    err_sys("waitpid error");

/*
 * We're the parent (the original process); we continue executing,
 * knowing that we're not the parent of the second child.
 */
exit(0);
}
```

Execution

```
$ ./a.out
$ second child, parent pid = 1
```

exec()

```
#include <unistd.h>
```

```
int execl(const char *pathname, const char *arg0, ... /* (char *)0 */);
```

```
int execlv(const char *pathname, char *const argv []);
```

```
int execlp(const char *pathname, const char *arg0, ... /* (char *)0,  
            char *const envp[] */);
```

```
int execve(const char *pathname, char *const argv[], char *const envp []);
```

```
int execlp(const char *filename, const char *arg0, ... /* (char *)0 */);
```

```
int execlvp(const char *filename, char *const argv []);
```

All six return: -1 on error, no return on success

Execute a program.

- When a process calls one of the exec functions, the process is completely replaced by the new program.
- New program starts executing at its main() function.
- exec() merely replaces the current process(text, data) with a brand new program from disk.

exec()

pathname(filename) argument

- pathname(filename) of a file to be executed.
- execl, execv, execl, execve take a pathname argument.
- execl_p, execv_p take a filename argument.
 - If filename contains a '/', it is taken as a pathname.
 - Otherwise, the executable file is searched for in the directories specified by PATH.
 - E.g. PATH=/bin:/usr/bin:/usr/local/bin/..

exec()

Argument list

- `arg0, arg1, ..., argn` in the `execl`, `execlp`, and `execle`.
- The list of arguments is terminated by a NULL pointer.
 - `execl("/bin/ls", "ls", "-al", 0);`

Argument vector

- `*argv[]` in `execv`, `execve` and `execvp`.
- The array of pointers is terminated by a NULL pointer.
 - `char *argv[3] = {"ls", "-al", 0};`
 - `execv("/bin/ls", argv);`

exec()

Environment variable

- *envp[] in `execv` and `execle`
- A pointer to an array of pointers to the environment strings.
- The other functions use the *environ* variable.
 - `char *env[2] = {"TERM=vt100", 0};`
 - `execle("/bin/ls", "ls", 0, env);`

exec()


Differences among the six exec functions

function	pathname	filename	arg list	argv[]	environ	envp[]
execl	0		0		0	
execlp		0	0		0	
execle	0		0			0
execv	0			0	0	
execvp		0		0	0	
execve	0			0		0
(letter in name)		p	l	v		e

exec()

❏ `execve` is a system call.

- `execl`, `execv`, `execle`, `execlp`, `execvp` are library functions.
- Relationship of the six exec functions.



exec()

Example

```
#include "apue.h"
#include <sys/wait.h>

char  *env_init[] = { "USER=unknown", "PATH=/tmp", NULL };

int
main(void)
{
    pid_t  pid;

    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid == 0) { /* specify pathname, specify environment */
        if (execle("/home/sar/bin/echoall", "echoall", "myarg1",
                  "MY ARG2", (char *)0, env_init) < 0)
            err_sys("execle error");
    }
}
```


exec()

Example(cont.)

```
if (waitpid(pid, NULL, 0) < 0)
    err_sys("wait error");

if ((pid = fork()) < 0) {
    err_sys("fork error");
} else if (pid == 0) { /* specify filename, inherit environment */
    if (execlp("echoall", "echoall", "only 1 arg", (char *)0) < 0)
        err_sys("execlp error");
}

exit(0);
}
```

exec()

echoall

```
#include "apue.h"

int
main(int argc, char *argv[])
{
    int    i;
    char    **ptr;
    extern char **environ;

    for (i = 0; i < argc; i++)           /* echo all command-line args */
        printf("argv[%d]: %s\n", i, argv[i]);

    for (ptr = environ; *ptr != 0; ptr++) /* and all env strings */
        printf("%s\n", *ptr);

    exit(0);
}
```

exec()

Execution

```
$ ./a.out
argv[0]: echoall
argv[1]: myarg1
argv[2]: MY ARG2
USER=unknown
PATH=/tmp
$ argv[0]: echoall
argv[1]: only 1 arg
USER=sar
LOGNAME=sar
SHELL=/bin/bash
...
HOME=/home/sar
```

system()

```
#include <stdlib.h>
```

```
int system(const char *cmdstring);
```

Returns: (see below)

 Consists of fork(), exec(), and waitpid().

- If fork() fails or waitpid() returns an error other than EINTR, returns -1.
- If exec() fails, return value is 127.
- If successful, return the termination status of the shell.

system()

An implementation of system()

```
#include <sys/wait.h>
#include <errno.h>
#include <unistd.h>

int
system(const char *cmdstring)
{
    pid_t pid;
    int status;

    if (cmdstring == NULL)
        return(1);

    if ((pid = fork()) < 0) {
        status = -1;
    } else if (pid == 0) {
        execl("/bin/sh", "sh", "-c", cmdstring, (char *)0);
        _exit(127);
    } else {
        /* probably out of processes */
        /* child */
        /* execl error */
    }
}
```

-c option: commands are read from 'cmdstring'

system()

An implementation of system() (cont.)

```
    } else {                                     /* parent */
        while (waitpid(pid, &status, 0) < 0) {
            if (errno != EINTR) {
                status = -1; /* error other than EINTR from waitpid() */
                break;
            }
        }
    }

    return(status);
}
```

system()

Example

```
#include "apue.h"
#include <sys/wait.h>
```

```
int main(void)
```

```
{
```

```
    int    status;
```

```
    if ((status = system("date")) < 0)
```

```
        err_sys("system() error");
```

```
    pr_exit(status);
```

pr_exit() is defined in “wait() and waitpid()” section.



```
    if ((status = system("nosuchcommand")) < 0)
```

```
        err_sys("system() error");
```

```
    pr_exit(status);
```

```
    if ((status = system("who; exit 44")) < 0)
```

```
        err_sys("system() error");
```

```
    pr_exit(status);
```

```
    exit(0);
```

```
}
```

system()

Execution

```
$ ./a.out
Sun Mar 21 18:41:32 EST 2004
normal termination, exit status = 0           for date
sh: nosuchcommand: command not found
normal termination, exit status = 127        for nosuchcommand
sar  :0   Mar 18 19:45
sar  pts/0 Mar 18 19:45 (:0)
sar  pts/1 Mar 18 19:45 (:0)
sar  pts/2 Mar 18 19:45 (:0)
sar  pts/3 Mar 18 19:45 (:0)
normal termination, exit status = 44        for exit
$
```